

Title of the use case	Description	Details
Use Case – 1 AI/ML-Driven Trade Confirmation Automation Objective – Automating Trade Confirmation creation using AI/ML technologies to reduce manual effort, improve accuracy, and accelerate turnaround times by leveraging intelligent template reuse, difference highlighting, and controlled generation of Confirmations.	Challenges – <ul style="list-style-type: none"> Manual Drafting & Operational Inefficiencies <ul style="list-style-type: none"> Redundancy in Drafting - Manual drafting leads to repetitive creation of similar templates, increasing redundancy and inefficiency in workflows. Time and Error Risks - The manual process is time-consuming and prone to human error, which raises operational risk and lowers productivity. Compliance and Standardization Issues - Inconsistent wording across documents undermines compliance and may cause regulatory exposure and legal issues. Need for Automation - Automation can improve template management, error reduction, and compliance, enabling scalability for global operations. Benefits – <ul style="list-style-type: none"> Time Savings: Reduce drafting effort by 60–80%. Quality & Compliance: Lower error rates; enforce regulatory norms. Scalability: Supports multiple trade types, counterparties, and jurisdictions. Adaptability: Extendable for multilingual confirmations and new products. Measurable Success Metrics - Key metrics include <ul style="list-style-type: none"> Precision ≥ 0.8, Clause Alignment F1 ≥ 0.85 Autofill Accuracy $\geq 95\%$. 	<p style="text-align: center;"><u>Template Reuse (Iteration 1)</u></p> <ul style="list-style-type: none"> Semantics Search for Templates – System embeddings-based retrieval to find and rank templates relevant to counterparty and trade type. Centralized Template Repository - A searchable, centralized repository stores templates to improve access and consistency in document drafting Enhanced Efficiency and Compliance - Eliminate blank-slate drafting; promotes consistency and improves compliance with standards. <p style="text-align: center;"><u>Difference Highlighting (Iteration 2)</u></p> <ul style="list-style-type: none"> Difference Highlighting Engine - to compare trade terms with templates to identify discrepancies in critical fields. Risk Severity Tagging – Tagging Differences (economic terms, legal clauses etc.) with risk indicators to prioritize high-risk changes for review. Combined Numeric and Semantic Analysis - The system uses numeric comparison and semantic analysis for comprehensive discrepancy detection. Integration and Compliance - Integration with clause libraries and rule-based validation ensures accuracy and compliance. <p style="text-align: center;"><u>Pre-Populated Template Generation (Iteration 3)</u></p> <ul style="list-style-type: none"> Automated Template Generation - The system uses AI and rule-based validation to pre-populate trade confirmation templates accurately and efficiently. Compliance and Controlled Output - Output is restricted to approved clauses ensuring compliance with regulatory requirements through deterministic validation. Productivity and Scalability Benefits - This iteration reduces manual effort, shortens cycle time, and scales across trade types and jurisdictions effectively. Future Enhancements - The solution enables future improvements like multilingual support and integration with downstream settlement systems.

Title of the use case	Description	Details
<p>Use Case 2 – AI Powered Knowledge Chatbot</p> <p>Objective –</p> <p>Develop a secure, low-cost, AI-powered knowledge chatbot to assist the internal operations teams in quickly finding policies, SOPs, product rules, exception handling steps, and compliance references—reducing resolution time and improving consistency.</p>	<p>Challenges –</p> <p>Current Challenges in Banking (Ops):</p> <ul style="list-style-type: none"> Information fragmentation: SOPs, product guides, and exception rules spread across SharePoint, Confluence, emails, PDFs, legacy portals. Policy churn: Frequent updates to regulatory and internal policies; hard to track latest versions. SME bottlenecks: High dependence on a few experts leads to delays. Manual search fatigue: Keyword search fails due to format/terminology mismatch (PDF scans, tables, appendices). Audit & compliance gaps: Limited traceability of what guidance was used, when, and by whom. Onboarding inefficiency: New joiners take weeks to learn where/how to find answers. <p>Benefits –</p> <p>Expected Impact (12–16 weeks MVP):</p> <ul style="list-style-type: none"> 30–50% reduction in time to find answers. 20–35% deflection of routine queries from SMEs. Faster onboarding of new operations staff (cut ramp-up time by 25–40%). Improved auditability and standardized responses. Improved ROI with Reduced Human intervention for L1 Support over period of time 	<p>The Solution –</p> <ul style="list-style-type: none"> Develop AI powered Banking Assistant Chatbot providing context-aware responses with Source Citations from company's existing Knowledge Base (viz. <i>SharePoint/Confluence sites, Network drives, emails, ticketing systems (ServiceNow/Jira), and bank policy libraries etc.</i>) Uses an open-source, Retrieval-Augmented Generation (RAG) architecture with <ul style="list-style-type: none"> on-prem orchestration, vector search, and role-based access, Query categorization: Users choice to choose specific domain/area or generic knowledge base for faster churning & response retrieval. Chat history management: The application maintains a history of users interactions to provide context aware responses Custom embeddings and retrieval - uses vector database for document retrieval and embeddings for query and document matching LLM integration: the chatbot integrates with a custom LLM generate responses (using Citi managed stellar API) based on the query and retrieved data Informative nature: No actions are taken automatically by the chatbot, it only provides information to the user enabling them to make effective decisions.